

The DataStax logo, featuring the word "DATASTAX" in a sans-serif font with a stylized cluster of blue dots to the right.

Introducing big data and Cassandra

Apache Cassandra:
Core Concepts, Skills, and Tools

Leo Schuman, Joe Chu

Oct 20, 2014

Learning Objectives

- **Understand big data**
- Describe Cassandra's history
- Understand common use cases
- Survey the Cassandra architecture

What are big data systems?


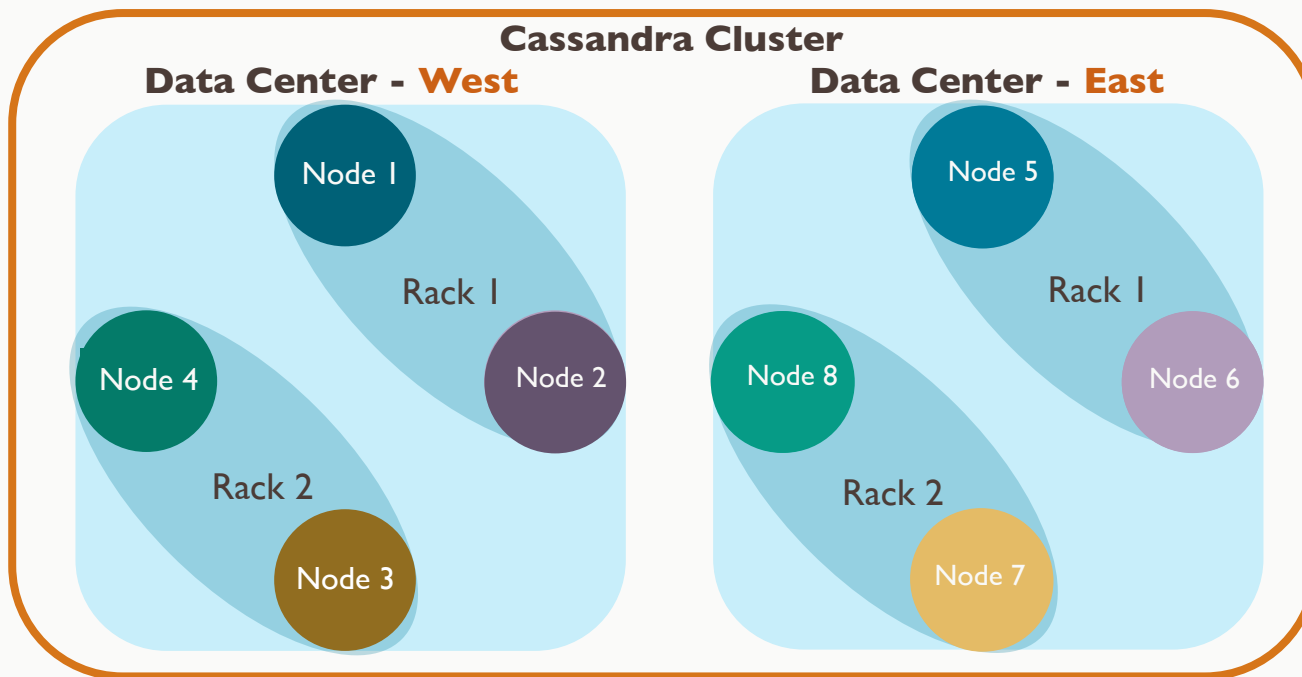
- Applications involving the "three V's"
 - **Volume**: gigabytes, growing to terabytes and beyond
 - **Velocity**: sensor data, click streams, financial transactions
 - **Variety**: data must be ingested from many different formats
- Characteristics requiring
 - multi-region availability
 - very fast and reliable response
 - no single point of failure

Why not relational data?

- Relational model provides
 - Normalized table schema
 - Cross table joins
 - ACID compliance
- But, at very high cost
 - Big data table joins – *billions* of rows, or more – require massive overhead
 - Sharding tables across systems is complex and fragile
- Modern applications have different priorities
 - Needs for speed and availability trump "always on" consistency
 - Commodity server racks trump massive high-end systems
 - Real world need for transactional guarantees is limited

What strategies help manage big data?

- Distribute data across nodes
- Relax consistency requirements
- Relax schema requirements
- Optimize data to suit actual needs

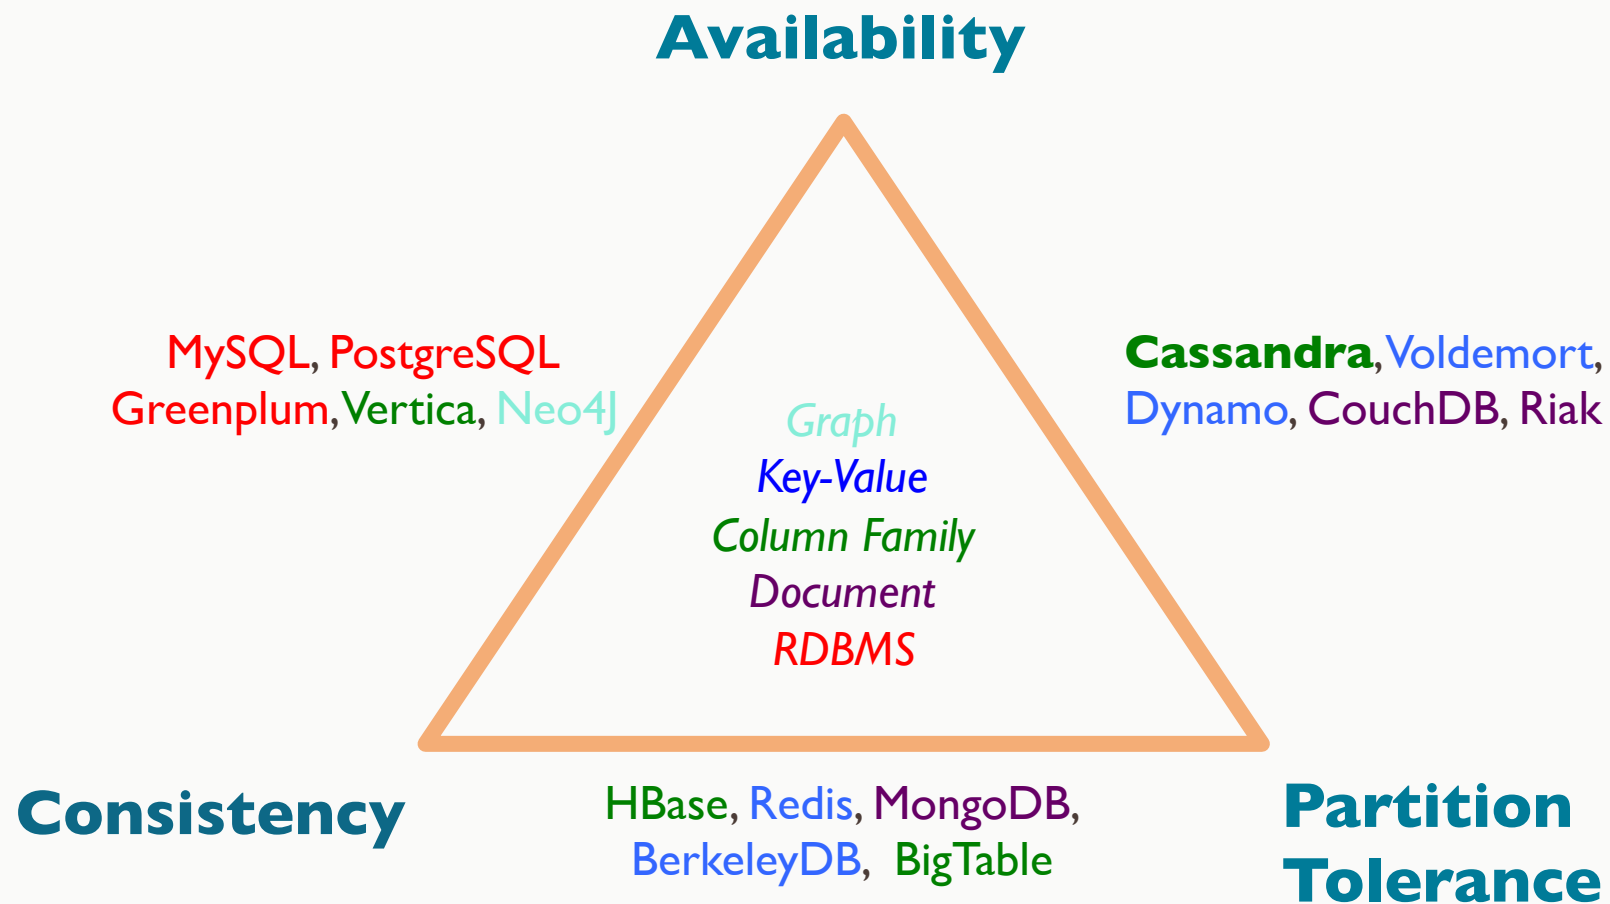


What is the NoSQL landscape?


- Four broad classes of non-relational database
 - **Graph**: data elements each relate to n others in a graph/network
 - **Key-Value**: keys map to arbitrary values of any data type
 - **Document**: document sets (JSON) queryable in whole or part
 - **Column Family**: keys mapped to sets of n -number of typed columns
- Three key factors help navigate the landscape
 - **Consistency**: do you get identical results, regardless which node is queried?
 - **Availability**: can the cluster respond to very high write and read volumes?
 - **Partition Tolerance**: is the cluster still available when part of it goes dark?

What is the CAP Theorem?

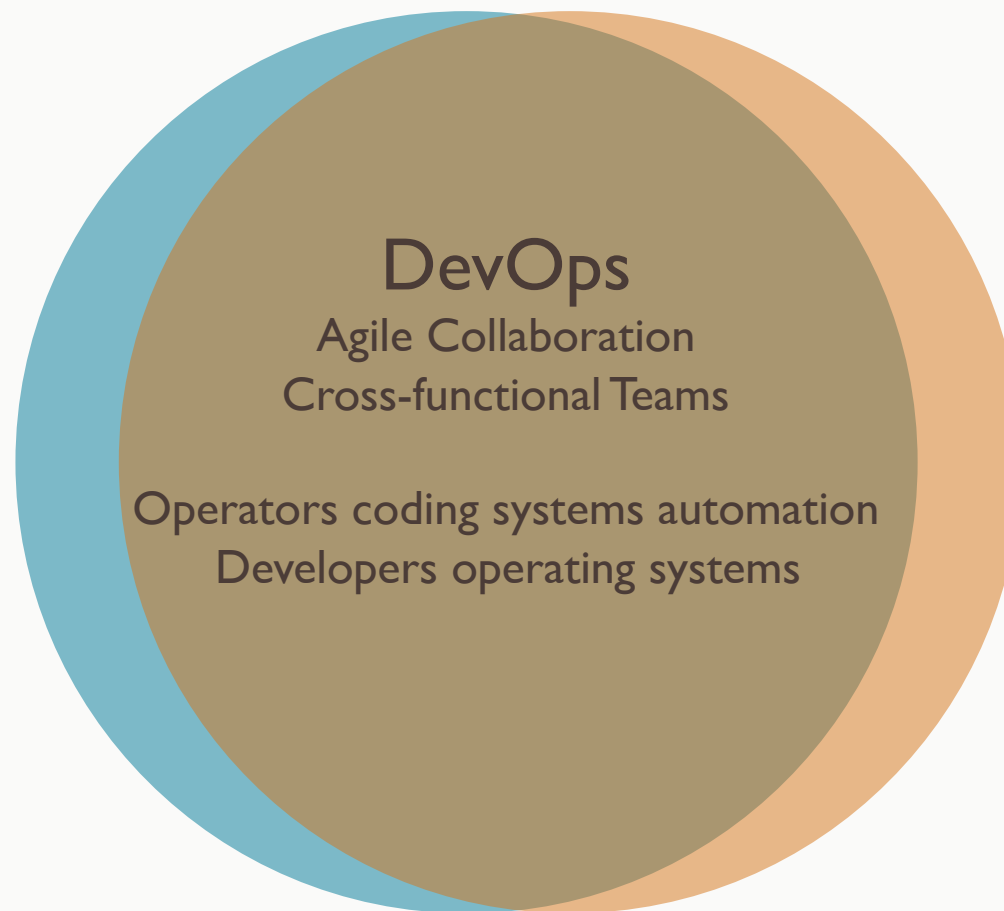
- In distributed systems, *consistency*, *availability*, and *partition tolerance* exist in a mutually dependent relationship. Pick any two.



What does "DevOps" mean?



What does "DevOps" mean?




Learning Objectives

- Understand big data
- **Describe Cassandra**
- Understand common use cases
- Survey the Cassandra architecture

What is Cassandra?

- Massively linearly scalable NoSQL database
 - Fully distributed, with no single point of failure
 - Free and open source, with deep developer support
 - Highly performant, with near-linear *horizontal* scaling in proper use cases



What is Cassandra?

- No single point of failure, due to horizontal scaling
 - *horizontal scaling*: add commodity hardware to a cluster
 - *vertical scaling*: add RAM and CPUs to a specialized high performance box



How has Cassandra evolved?

- Core technologies
 - Google BigTable: foundation of the storage model
 - Amazon Dynamo: foundation of the distribution backbone
 - Facebook: integrated BigTable and Dynamo, then released as Cassandra
- Rapid evolution
 - 0.6 – April 2010
 - 0.7 – January 2011
 - 0.8 – June 2011
 - 1.0 – October 2011
 - 1.1 – April 2012
 - 1.2 – January 2013
 - 2.0 – September 2013
 - 2.1 – September 2014
- Top-level Apache project since 2010

How does Cassandra model data?

- **Cassandra Query Language (CQL)**
 - Provides a familiar, row-column, SQL-like approach
 - CREATE, ALTER, DROP
 - SELECT, INSERT, UPDATE, DELETE
 - Replaced the complex, storage-oriented Thrift API used in prior versions
 - Provides clear schema definitions in a flexible (NoSQL) schema context

```
CREATE TABLE Performer (  
    name VARCHAR,  
    type VARCHAR,  
    country VARCHAR,  
    style VARCHAR,  
    born INT,  
    died INT,  
    PRIMARY KEY (name)  
);
```


Learning Objectives

- Understand big data
- Describe Cassandra's history
- **Understand common use cases**
- Survey the Cassandra architecture

When is Cassandra the best solution?

- Cassandra excels when you need
 - No single point of failure
 - Real-time writes with live operational data analysis
 - Flexible, easily altered data models
 - Near-linear horizontal scaling across commodity servers
 - Reliable replication across distributed data centers
 - Clearly defined table schema in a NoSQL environment


When is Cassandra not the best solution?

- Traditional RDBMS excels when you need
 - ACID-compliant transactions, with rollback (e.g., bank transfers)
 - Justification for high-end hardware

What are common Cassandra use cases?

- Cassandra is particularly useful for
 - Playlists and collections (such as Spotify)
 - Personalization and recommendation engines (such as Ebay)
 - Messaging (such as Instagram)
 - Fraud detection (such as Barracuda)
 - Sensor data (such as Zonar)
- Many, many functional and industry use cases available
 - <http://planetcassandra.org/functional-use-cases/>

Who is using Cassandra?



*for more, see
[PlanetCassandra.org/
companies/](http://PlanetCassandra.org/companies/)*

Summary

- "Big data" requires very high availability, tolerance, and response
- Manage with relaxed consistency and optimization to actual need
- There are four categories of NoSQL systems: graph, key-value, document, and column family
- Cassandra is a column-family system, offering near-linear scaling across commodity clusters
- Developer and Operator roles are merging into "DevOps"
- Cassandra has its roots in Amazon Dynamo and Google BigTable technology
- Data is modeled using *Cassandra Query Language (CQL)*
- Cassandra excels with flexible, real-time data ingestion and analysis with no single point of failure across commodity hardware clusters

Review Questions

- What do consistency, availability, and partition tolerance mean?
- Where does Cassandra fit within the CAP theorem?
- What are the technological roots of Cassandra?
- What technology does Cassandra use to model data?
- What use cases are great fits for Cassandra?

