



# **Apache Cassandra:** Core Concepts, Skills, and Tools

**Introducing the Cassandra Data Model  
and Cassandra Query Language**  
Exercise Workbook

Artem Chebotko  
Leo Schuman  
October, 2014

## Exercise I: Model sample data as column families

### In this exercise, you will:

- Explore sample user data
- Organize data into column families

### Steps

#### Explore sample user data


1. Shown below, understand the tabular data that describes users.

Name	DOB	Email	Join Date
John	12/01/1986	john@data.org	03/08/2014
Mary	12/01/1986	mary@data.org	03/08/2014
John	02/18/1979	john@data.edu	01/01/2013

2. In the space below or a separate sheet, list the fields above which uniquely identify a user.


## Organize data into column families

3. In the template below or a separate sheet, fill in blanks to show how a column family can store user data in rows, where each row is identified by an email.




## Introducing the Cassandra Data Model and Cassandra Query Language

4. In the template below or a separate sheet, fill in blanks to show how a column family can store user data in rows, where each row is identified by a name and a date of birth.



5. In the template below or a separate sheet, fill in blanks to show how a column family can store user data in rows, where each row is identified by a join date.



END OF EXERCISE

## Exercise 2: Represent column families as tables


**In this exercise, you will:**

- Explore sample column families
- Represent column families as tables

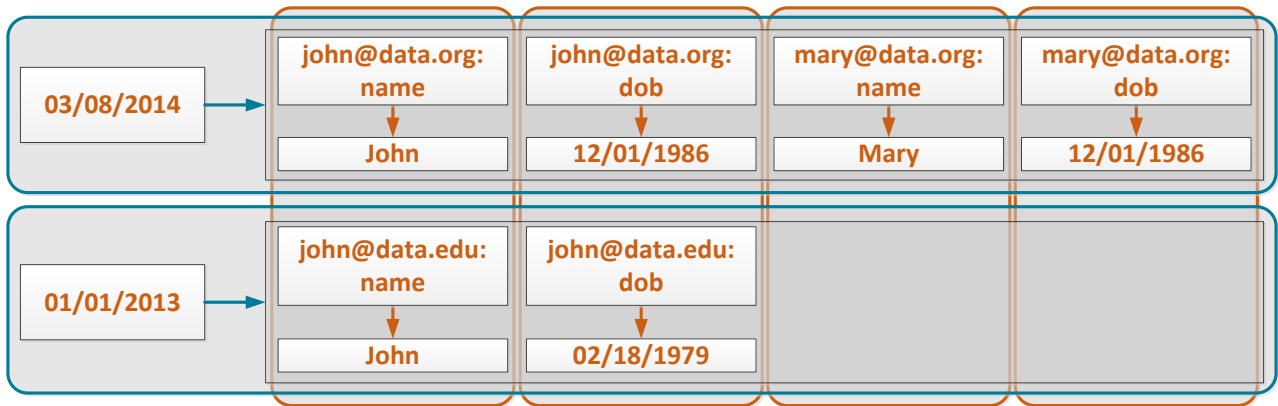
### Steps

#### Explore sample column families

1. Shown below, consider sample column families that describe users.



## Introducing the Cassandra Data Model and Cassandra Query Language



2. In the space below or a separate sheet, list a row key (simple or composite) and column keys (simple or composite) for each column family.

## Represent column families as tables

3. In the templates below or a separate sheet, show how user data can be stored in tables.

## users\_by\_email


## users\_by\_name\_and\_dob


users\_by\_join\_date


4. In the space below or a separate sheet, list a partition key (simple or composite), clustering columns, and a primary key for each table.



END OF EXERCISE

## Demo 3: How to launch and use *cqlsh*

### In this demo, we will:

- Show how to start *cqlsh*
- Show how to execute a CQL script using *cqlsh*
- Show how to explore a database schema and instance

### Steps

#### Show how to start *cqlsh*

1. In a terminal window, start *cqlsh*.

```
ccm node1 cqlsh
```

*This demonstration can be run in either the cascor cluster or tarball installation.*

#### Show how to execute a CQL script using *cqlsh*

2. In *cqlsh*, using the *SOURCE* command, execute script *userdb.cql* for this demo, located at *~/cascor/intro-dm-cql/demo-3*. The script content is shown below.

```
SOURCE '~/cascor/intro-dm-cql/demo-3/userdb.cql';
```

```
CREATE KEYSPACE userdb
WITH replication = {'class': 'SimpleStrategy',
'replication_factor' : 1};

USE userdb;

CREATE TABLE users_by_email (
  name VARCHAR,
  dob TIMESTAMP,
  email VARCHAR,
  join_date TIMESTAMP,
  PRIMARY KEY (email)
);
```



```
CREATE TABLE users_by_name_and_dob (  
    name VARCHAR,  
    dob TIMESTAMP,  
    email VARCHAR,  
    join_date TIMESTAMP,  
    PRIMARY KEY ((name,dob))  
);  
  
CREATE TABLE users_by_join_date (  
    name VARCHAR,  
    dob TIMESTAMP,  
    email VARCHAR,  
    join_date TIMESTAMP,  
    PRIMARY KEY (join_date, email)  
);  
  
INSERT INTO users_by_email (name, dob, email, join_date)  
VALUES ('John', '1986-12-01', 'john@data.org', '2014-03-08');  
INSERT INTO users_by_email (name, dob, email, join_date)  
VALUES ('Mary', '1986-12-01', 'mary@data.org', '2014-03-08');  
INSERT INTO users_by_email (name, dob, email, join_date)  
VALUES ('John', '1979-02-18', 'john@data.edu', '2013-01-01');  
  
INSERT INTO users_by_name_and_dob (name, dob, email,  
join_date) VALUES ('John', '1986-12-01', 'john@data.org',  
'2014-03-08');  
INSERT INTO users_by_name_and_dob (name, dob, email,  
join_date) VALUES ('Mary', '1986-12-01', 'mary@data.org',  
'2014-03-08');  
INSERT INTO users_by_name_and_dob (name, dob, email,  
join_date) VALUES ('John', '1979-02-18', 'john@data.edu',  
'2013-01-01');  
  
INSERT INTO users_by_join_date (name, dob, email,  
join_date) VALUES ('John', '1986-12-01', 'john@data.org',  
'2014-03-08');  
INSERT INTO users_by_join_date (name, dob, email,  
join_date) VALUES ('Mary', '1986-12-01', 'mary@data.org',  
'2014-03-08');  
INSERT INTO users_by_join_date (name, dob, email,  
join_date) VALUES ('John', '1979-02-18', 'john@data.edu',  
'2013-01-01');
```

**Show how to explore a database schema and instance**

3. In *cqlsh*, run *USE* and *DESCRIBE KEYSPACE* commands for *userdb*.
4. In *cqlsh*, run *DESCRIBE TABLES* and *DESCRIBE TABLE* commands for each table.
5. In *cqlsh*, run *SELECT* statement to retrieve all rows from each table.
6. In *cqlsh*, run *EXIT* command.

END OF DEMO

## Exercise 4: Create a keyspace and tables using *cqlsh*

### In this exercise, you will:

- Create a keyspace
- Create tables in a keyspace
- Populate tables from CSV files
- Execute simple queries

### Steps

#### Create a keyspace

1. In the terminal window, navigate to the *cascor/intro-dm-cql* directory.

```
cd ~/cascor/intro-dm-cql
```

2. From the *cascor/intro-dm-cql* directory, start up *cqlsh*.

```
ccm node1 cqlsh
```

3. In *cqlsh*, create a new keyspace named *musicdb* for a music database, with the simple replication strategy, and a replication factor of 3.

```
CREATE KEYSPACE musicdb
WITH replication = {
  'class': 'SimpleStrategy',
  'replication_factor' : 3
};
```

4. In *cqlsh*, use the DESCRIBE command to display information about the keyspace.

```
DESCRIBE KEYSPACE musicdb
```

## Create tables in a keyspace

5. In *cqlsh*, use the `USE` command to set the *musicdb* keyspace as the current default.

```
USE musicdb;
```

6. In *cqlsh*, create the tables for the *musicdb* keyspace by executing the CQL script.

```
SOURCE '~/cascor/intro-dm-cql/exercise-4/musicdb.cql'
```

7. Shown below, study and understand the table definitions that were created from the CQL script.

```
CREATE TABLE performer (  
  name VARCHAR,  
  type VARCHAR,  
  country VARCHAR,  
  style VARCHAR,  
  founded INT,  
  born INT,  
  died INT,  
  PRIMARY KEY (name)  
);  
  
CREATE TABLE performers_by_style (  
  style VARCHAR,  
  name VARCHAR,  
  PRIMARY KEY (style, name)  
);
```

```
CREATE TABLE album (  
    title VARCHAR,  
    year INT,  
    performer VARCHAR,  
    genre VARCHAR,  
    tracks MAP<INT,VARCHAR>,  
    PRIMARY KEY ((title, year))  
);  
  
CREATE TABLE albums_by_performer (  
    performer VARCHAR,  
    year INT,  
    title VARCHAR,  
    genre VARCHAR,  
    PRIMARY KEY (performer, year, title)  
) WITH CLUSTERING ORDER BY (year DESC, title ASC);  
  
CREATE TABLE albums_by_genre (  
    genre VARCHAR,  
    performer VARCHAR,  
    year INT,  
    title VARCHAR,  
    PRIMARY KEY (genre, performer, year, title)  
) WITH CLUSTERING ORDER BY (performer ASC, year DESC,  
title ASC);  
  
CREATE TABLE albums_by_track (  
    track_title VARCHAR,  
    performer VARCHAR,  
    year INT,  
    album_title VARCHAR,  
    PRIMARY KEY (track_title, performer, year, album_title)  
) WITH CLUSTERING ORDER BY (performer ASC, year DESC,  
album_title ASC);  
  
CREATE TABLE tracks_by_album (  
    album_title VARCHAR,  
    year INT,  
    performer VARCHAR STATIC,  
    genre VARCHAR STATIC,  
    number INT,  
    track_title VARCHAR,  
    PRIMARY KEY ((album_title, year), number)  
);
```

8. In *cqlsh*, display the names of all tables in the *musicdb* keyspace.

### DESCRIBE TABLES

### Populate tables from CSV files

9. In *cqlsh*, run the commands below one at a time to import data into the tables from CSV files.

*You may run into an intermittent bug that causes the COPY command to abort or to crash cqlsh, as reported in the JIRA ticket CASSANDRA-8351. This may show errors such as:*

line contains NULL byte  
Aborting import at record #0. Previously-inserted values still present.

Segmentation fault (core dumped)

*Retry the previous command as necessary until it succeeds.*

```
COPY performer
  (name, type, country, style, founded, born, died)
FROM '~/cascor/intro-dm-cql/exercise-4/performer.csv'
WITH HEADER = 'true';

COPY performers_by_style (style, name)
FROM '~/cascor/intro-dm-cql/exercise-4/performers_by_style.csv'
WITH HEADER = 'true';

COPY album (title, year, performer, genre, tracks)
FROM '~/cascor/intro-dm-cql/exercise-4/album.csv'
WITH HEADER = 'true';

COPY albums_by_performer (performer, year, title, genre)
FROM '~/cascor/intro-dm-cql/exercise-4/albums_by_performer.csv'
WITH HEADER = 'true';

COPY albums_by_genre (genre, performer, year, title)
FROM '~/cascor/intro-dm-cql/exercise-4/albums_by_genre.csv'
WITH HEADER = 'true';

COPY albums_by_track (track_title, performer, year, album_title)
FROM '~/cascor/intro-dm-cql/exercise-4/albums_by_track.csv'
WITH HEADER = 'true';

COPY tracks_by_album
  (album_title, year, performer, genre, number, track_title)
FROM '~/cascor/intro-dm-cql/exercise-4/tracks_by_album.csv'
WITH HEADER = 'true';
```

## Execute simple queries

10. In *cqlsh*, test the following queries.

```
SELECT * FROM performer WHERE name = 'The Beatles';
```

```
SELECT * FROM performer WHERE name = 'John Lennon';
```

*Notice the differences in data that is stored for the band and the artist.*

11. In *cqlsh*, test the following queries.

```
SELECT *  
FROM tracks_by_album  
WHERE album_title = 'Revolver' AND year = 1966;
```

```
SELECT *  
FROM album  
WHERE title = 'Revolver' AND year = 1966;
```

*Notice how the same data is organized differently in these tables.*

END OF EXERCISE

## Exercise 5: Create tables using UUID, TIMEUUID, and COUNTER columns

**In this exercise, you will:**

- Select an existing keyspace
- Create additional tables with UUID, TIMEUUID, and COUNTER columns

### Steps

#### Select an existing keyspace

1. In *cqlsh*, use the `USE` command to set the *musicdb* keyspace as the current default.

```
USE musicdb;
```

*Note that you may also use the keyspace name as a table name prefix. The `USE` command, however, adds convenience by setting a default keyspace so that no prefix is required.*

#### Create additional tables with UUID, TIMEUUID, and COUNTER columns

2. In *cqlsh*, create table *user*.

```
CREATE TABLE user (  
    id UUID,  
    name VARCHAR,  
    PRIMARY KEY (id)  
);
```



3. In *cqlsh*, create table *track\_ratings\_by\_user*.

```
CREATE TABLE track_ratings_by_user (  
  user UUID,  
  activity TIMEUUID,  
  rating INT,  
  album_title VARCHAR,  
  album_year INT,  
  track_title VARCHAR,  
  PRIMARY KEY (user, activity)  
) WITH CLUSTERING ORDER BY (activity DESC);
```

*Note that the rating activities for a user will be sorted in the descending order of their timestamps (most recent first) extracted from the TIMEUUID column.*

4. In *cqlsh*, create table *ratings\_by\_track*.

```
CREATE TABLE ratings_by_track (  
  album_title VARCHAR,  
  album_year INT,  
  track_title VARCHAR,  
  num_ratings COUNTER,  
  sum_ratings COUNTER,  
  PRIMARY KEY (album_title, album_year, track_title)  
);
```

*Note that all non-counter columns are part of the primary key.  
Also, note that an average rating can be computed as  
 $avg\_rating = sum\_ratings / num\_ratings$ .*

*We will populate these tables in a different exercise.*

END OF EXERCISE

## Exercise 6: Add user-defined type, alter tables, add collection column, and add secondary indexes

### In this exercise, you will:

- Create a user-defined type
- Add and drop columns to an existing table
- Add secondary indexes to an existing table

### Steps

#### Create a user-defined type

1. In *cqlsh*, set the *musicdb* keyspace as the current default.

```
USE musicdb;
```

2. In *cqlsh*, display information about available user-defined types in *musicdb*.

```
DESCRIBE TYPES
```

3. In *cqlsh*, create a new user-defined type called *track*.

```
CREATE TYPE track (  
    album_title VARCHAR,  
    album_year INT,  
    track_title VARCHAR  
);
```

4. In *cqlsh*, display information about types.

```
DESCRIBE TYPES
```

5. In *cqlsh*, display information about the user-defined type *track*.

```
DESCRIBE TYPE track
```

### Add columns to an existing table

6. In *cqlsh*, drop the columns *album\_title*, *album\_year*, and *track\_title* for the table *track\_ratings\_by\_user*.

```
ALTER TABLE track_ratings_by_user DROP album_title;  
ALTER TABLE track_ratings_by_user DROP album_year;  
ALTER TABLE track_ratings_by_user DROP track_title;
```

7. In *cqlsh*, add a new column called *song* to the *track\_ratings\_by\_user* table with the user-defined type *track*.

```
ALTER TABLE track_ratings_by_user ADD song frozen<track>;
```

8. In *cqlsh*, display information about the table *track\_ratings\_by\_user*.

```
DESCRIBE TABLE track_ratings_by_user
```

9. In *cqlsh*, display information about table *user*.

```
DESCRIBE TABLE user
```

10. In *cqlsh*, add two new columns to table *user*.

```
ALTER TABLE user ADD email VARCHAR;
```

```
ALTER TABLE user ADD preferences SET<VARCHAR>;
```

11. In *cqlsh*, display information about table *user*.

```
DESCRIBE TABLE user
```

## Add secondary indexes to an existing table

12. In *cqlsh*, create a secondary index on the *preferences* column in the table *user*.

```
CREATE INDEX user_preferences_key ON user (preferences);
```

13. In *cqlsh*, display information about table *user*.

```
DESCRIBE TABLE user
```

14. In *cqlsh*, display information about table *performer*.

```
DESCRIBE TABLE performer
```

15. In *cqlsh*, create two secondary indexes for table *performer*.

```
CREATE INDEX performer_country_key ON performer (country);
```

```
CREATE INDEX performer_style_key ON performer (style);
```

16. In *cqlsh*, display information about table *performer*.

```
DESCRIBE TABLE performer
```

17. In *cqlsh*, test the following query.

```
SELECT name  
FROM performer  
WHERE country = 'Iceland' AND style = 'Rock'  
ALLOW FILTERING;
```

END OF EXERCISE

## Demo 7: How to launch and use DevCenter

### In this demo, we will:

- Show how to start DevCenter
- Show how to connect to a Cassandra cluster and explore database objects
- Show how to create, edit, and execute a new CQL script


### Steps

#### Show how to start DevCenter

1. In a terminal window, start *DevCenter*.

#### Show how to connect to a Cassandra cluster

2. In *DevCenter*, using the *Connection Manager*, connect to the Cassandra cluster.



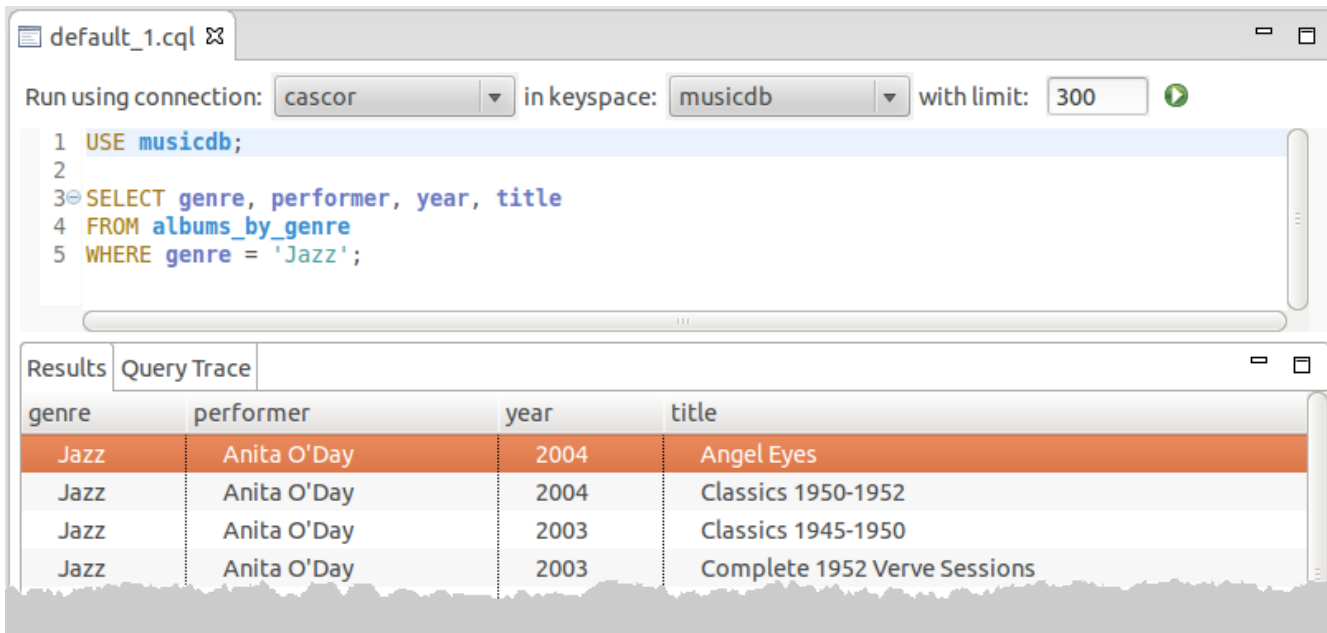
### Show how to explore existing database objects

3. In *DevCenter*, using the *Schema Explorer*, explore database objects in the *musicdb* keyspace.

### Show how to create, edit, and execute a new CQL script

4. In *DevCenter*, create a new CQL script, add *USE* and *SELECT* statements, and execute the script.

```
USE musicdb;  
SELECT genre, performer, year, title FROM albums_by_genre  
WHERE genre = 'Jazz';
```



default\_1.cql

Run using connection: **cascor** in keyspace: **musicdb** with limit: **300**

```
1 USE musicdb;  
2  
3 SELECT genre, performer, year, title  
4 FROM albums_by_genre  
5 WHERE genre = 'Jazz';
```

genre	performer	year	title
Jazz	Anita O'Day	2004	Angel Eyes
Jazz	Anita O'Day	2004	Classics 1950-1952
Jazz	Anita O'Day	2003	Classics 1945-1950
Jazz	Anita O'Day	2003	Complete 1952 Verve Sessions

#### Useful DevCenter shortcuts:

<Ctrl>+<s> to save a current CQL script  
<Alt>+<f>, then <.> to open an existing CQL script  
<Ctrl>+<space> to bring up the autocomplete menu  
<Shift>+<Ctrl>+<l> to open the key assist / shortcut menu  
<Alt>+<F1> to execute the current CQL script

END OF DEMO

## Exercise 8: Inserting and updating values using DevCenter

### In this exercise, you will:

- Insert rows into tables
- Update rows in tables

### Steps

#### Insert rows into tables

1. In *DevCenter*, create a new script with the following content.

```
USE musicdb;

INSERT INTO user (id, name, email) VALUES (12345678-abcd-
abcd-abcd-abcd12345678, 'John', 'john@datastax.com');

INSERT INTO user (id, name, email) VALUES (87654321-abcd-
abcd-abcd-abcd87654321, 'Mary', 'mary@datastax.com');

INSERT INTO user (id, name, email) VALUES (77777777-beef-
beef-beef-beef77777777, 'Joe', 'joe@datastax.com');
```

*Note that more readable UUIDs were selected on purpose.*

2. In *DevCenter*, execute the above script exactly one time.

*Make a note of the execution time for the three insert operations. You'll be comparing them with some other inserts in just a few steps.*

3. In *DevCenter*, create a new script with the query and execute it.

```
USE musicdb;  
  
INSERT INTO user (id, name, email) VALUES (00000000-aaaa-  
aaaa-aaaa-aaaa00000000, 'Ron', 'ron@datastax.com')  
IF NOT EXISTS;  
  
INSERT INTO user (id, name, email) VALUES (12345678-abcd-  
abcd-abcd-abcd12345678, 'Steve', 'steve@datastax.com')  
IF NOT EXISTS;
```

*Do these two INSERTs execute successfully? What is the value of the [applied] column? How long did these INSERT operations take to execute compared to the earlier ones?*

4. In *DevCenter*, create a new script with the query and execute it.

```
USE musicdb;  
  
SELECT * FROM user;
```

### Update rows in tables

5. In *DevCenter*, create a new script with the following content and execute it.

```
USE musicdb;  
  
SELECT * FROM performer WHERE name = 'The Beatles';
```

*Notice the year the band was founded.*

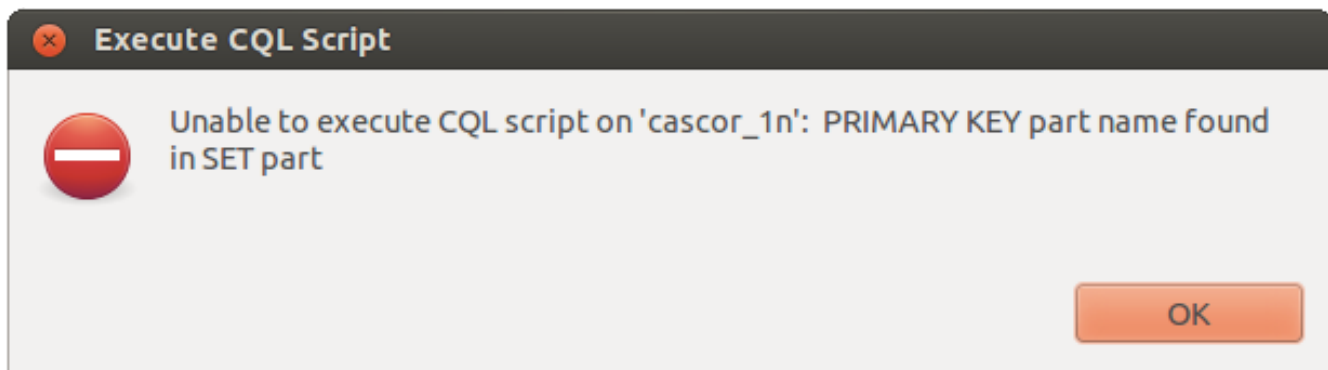
6. In *DevCenter*, modify the script to add the following statements and execute it.

```
USE musicdb;  
  
UPDATE performer  
SET founded = 1960  
WHERE name = 'The Beatles';  
  
SELECT * FROM performer WHERE name = 'The Beatles';
```



7. In DevCenter, modify the script with the following *UPDATE* statement and execute it.

```
USE musicdb;  
  
UPDATE performer  
SET name = 'Beatles'  
WHERE name = 'The Beatles';  
  
SELECT * FROM performer WHERE name = 'Beatles';
```



*Values in primary key columns cannot be updated. If such values change, an old row should be deleted and a new one should be inserted.*

8. In DevCenter, modify the script with the following *UPDATE* statement and execute it.

```
USE musicdb;  
  
UPDATE performer  
SET founded = 1957  
WHERE name = 'The Beatles'  
IF type = 'band';  
  
SELECT * FROM performer WHERE name = 'The Beatles';
```

*The UPDATE operation is only executed if the value for the type column is the same. What is the value of the returned [applied] column for this update?*

9. In *DevCenter*, modify the script with the following *UPDATE* statement and execute it.

```
USE musicdb;  
  
UPDATE performer  
SET born = 1960  
WHERE name = 'The Beatles'  
IF type = 'artist';  
  
SELECT * FROM performer WHERE name = 'The Beatles';
```

*Why did this update fail? What are the returned values for the [applied] and type column?*

END OF EXERCISE

## Exercise 9: Manipulate values in counter, collection and UDT columns

### In this exercise, you will:

- Update a counter column
- Update a collection column
- Update a UDT column

### Steps

#### Update a counter column

1. In *cqlsh*, select the *musicdb* keyspace to work with.

```
USE musicdb;
```

2. Display information about table *ratings\_by\_track*.

```
DESCRIBE TABLE ratings_by_track;
```

3. Update the counter columns.

```
UPDATE ratings_by_track
SET num_ratings = num_ratings + 1,
    sum_ratings = sum_ratings + 5
WHERE album_title = 'Revolver' AND
      album_year = 1966 AND
      track_title = 'Yellow Submarine';
```

4. Run the following query and observe the results.

```
SELECT * FROM ratings_by_track;
```

5. Update the counter columns again.

```
UPDATE ratings_by_track
SET num_ratings = num_ratings + 1,
    sum_ratings = sum_ratings + 4
WHERE album_title = 'Revolver' AND
      album_year = 1966 AND
      track_title = 'Yellow Submarine';
```

6. Run the following query and observe the results.

```
SELECT * FROM ratings_by_track;
```

### Update a collection column

7. In *cqlsh*, display information about table *user*.

```
DESCRIBE TABLE user;
```

8. Run the following query and observe the results.

```
SELECT * FROM user;
```

9. Update the set column for a specified user.

```
UPDATE user SET preferences = preferences +
                                     {'Rock', 'Classic'}
WHERE id = 12345678-abcd-abcd-abcd-abcd12345678;
```

10. Run the following query and observe the results.

```
SELECT * FROM user;
```

11. Update the set column for a specified user.

```
UPDATE user SET preferences = preferences +  
                                {'Rock', 'Jazz'}  
WHERE id = 12345678-abcd-abcd-abcd-abcd12345678;
```

12. In *cqlsh*, test the following query and observe the results.

```
SELECT * FROM user;
```

*Notice how the set values are ordered. Duplicate values have been eliminated.*

### Update a UDT column

13. Display information about table *track\_ratings\_by\_user*.

```
DESCRIBE TABLE track_ratings_by_user;
```

14. In *cqlsh*, insert a new rating for a specified user.

```
INSERT INTO track_ratings_by_user  
(user, activity, rating, song)  
VALUES(12345678-abcd-abcd-abcd-abcd12345678, 1234abcd-  
1234-1134-1234-abcd1234abcd, 8, {'album_title': 'what A  
wonderful world', album_year: 1968, track_title: 'Boogie  
After Midnight'});
```

15. In *cqlsh*, test the following query and observe the results.

```
SELECT * FROM track_ratings_by_user;
```

16. In *cqlsh*, update the *track\_title* to a new specified value.

```
UPDATE track_ratings_by_user SET song =  
{track_title: 'Mousetrap'}  
WHERE user = 12345678-abcd-abcd-abcd-abcd12345678 AND  
activity = 1234abcd-1234-1134-1234-abcd1234abcd;
```

17. In *cqlsh*, test the following query and observe the results.

```
SELECT * FROM track_ratings_by_user;
```

*Although the intention was just to update the track title, what happened to the album title and the album year?*

18. In *cqlsh*, update the *song* column to a new specified value.

```
UPDATE track_ratings_by_user SET song =  
{album_title: 'what A wonderful world', album_year: 1968,  
track_title: 'Mousetrap' }  
WHERE user = 12345678-abcd-abcd-abcd-abcd12345678 AND  
activity = 1234abcd-1234-1134-1234-abcd1234abcd;
```

19. In *cqlsh*, test the following query and observe the results.

```
SELECT * FROM track_ratings_by_user;
```

*Since the UDT column is frozen, all of the fields for the song column must be updated together.*

END OF EXERCISE

## Exercise 10: Explore equality and range search in queries

### In this exercise, you will:

- Understand and execute a set of queries

### Steps

#### Understand and execute a set of queries

1. In *DevCenter*, run the following query.

```
SELECT * FROM albums_by_track
WHERE track_title IN ('Yesterday', 'Tomorrow');
```

2. Run the following query.

```
SELECT * FROM tracks_by_album
WHERE album_title = '20 Greatest Hits' AND year = 1982 AND
      number > 8 AND number < 12;
```

3. Run the following query.

```
SELECT * FROM albums_by_performer
WHERE performer = 'The Beatles';
```

4. Run the following query.

```
SELECT * FROM albums_by_performer
WHERE performer = 'The Beatles' AND
      year >= 1960 AND year <= 1980;
```

5. Run the following query.

```
SELECT * FROM albums_by_performer
WHERE performer = 'The Beatles' AND
      year >= 1960 AND year <= 1980
ORDER BY year ASC;
```

6. Run the following query.

```
SELECT * FROM albums_by_performer
WHERE year >= 1960 AND year <= 1980
LIMIT 50 ALLOW FILTERING;
```

7. Run the following query.

```
SELECT * FROM performer
WHERE style = 'Rock';
```

8. Run the following query.

```
SELECT * FROM performer
WHERE style = 'Rock' AND country >= 'U'
LIMIT 1000
ALLOW FILTERING;
```

9. Run the following query.

```
SELECT * FROM user
WHERE preferences CONTAINS 'Rock';
```

END OF EXERCISE



## Demo I I: Explore queries with various predicates (optional)

### In this demo, we will:

- Showcase equality queries
- Showcase range queries

### Steps

#### Showcase equality queries

1. In *DevCenter*, design and execute different equality queries over an existing table (below is an example).

```
SELECT * FROM albums_by_genre;

SELECT * FROM albums_by_genre
WHERE genre = 'Rock';

SELECT * FROM albums_by_genre
WHERE genre IN ('Country', 'Rock');

SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND performer = 'The Beatles';

SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND
      performer = 'The Beatles' AND
      year = 1964;

SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND
      performer = 'The Beatles' AND
      year = 1964 AND
      title = 'Beatles For Sale';
```

## Showcase range queries

2. In *DevCenter*, design and execute different range queries over an existing table (below is an example).

```
SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND performer > 'T';

SELECT * FROM albums_by_genre
WHERE genre IN ('Country', 'Rock') AND
      performer > 'T' AND performer <= 'U2';

SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND
      performer = 'The Beatles' AND
      year < 1980;

SELECT * FROM albums_by_genre
WHERE genre = 'Rock' AND
      performer = 'The Beatles' AND
      year = 1964 AND
      title < 'C';
```

END OF DEMO

## Demo I 2: Explore online resources on CQL and data modeling

### In this demo, we will:

- Showcase online resources on CQL and data modeling
- Showcase educational video presentations on CQL and data modeling

### Steps

#### Showcase online resources on CQL and data modeling

1. In a web browser, go to <http://www.datastax.com/documentation> and explore existing documentation on data modeling and CQL. Point out CQL features that are not covered in this module.

#### Showcase educational video presentations on CQL and data modeling

2. In a web browser, go to <http://planetcassandra.org/summit-presentations/> and explore existing summit presentations, community webinars, meetup presentations, podcasts, and other available resources.

END OF DEMO