



Working with the Cassandra Read Path

Apache Cassandra:
Core Concepts, Skills, and Tools


Leo Schuman, Joe Chu

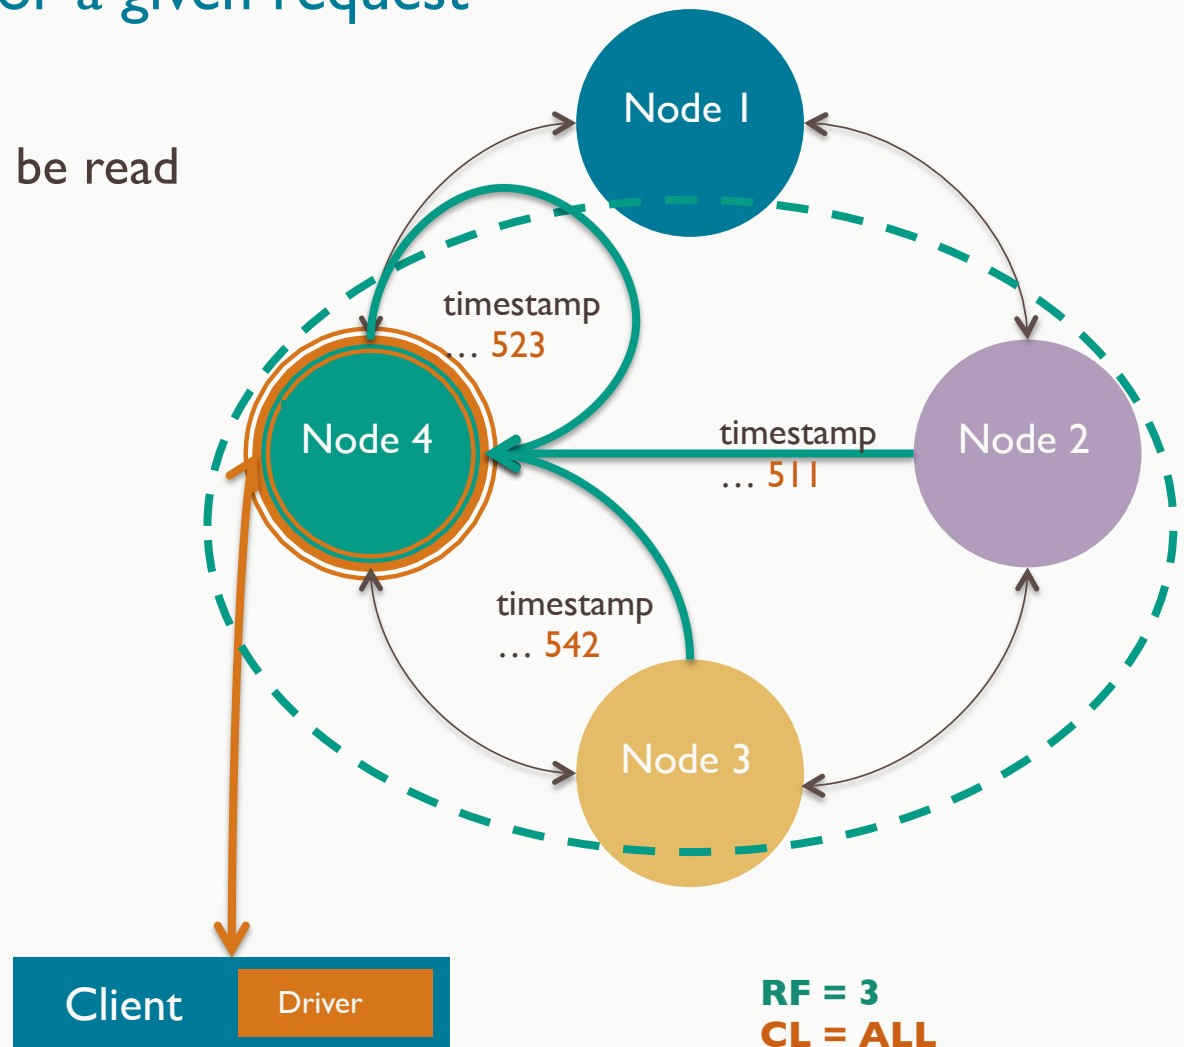
Oct 20, 2014

Learning Objectives

- **Understand how data is read from the storage engine**
- Read data from Cassandra

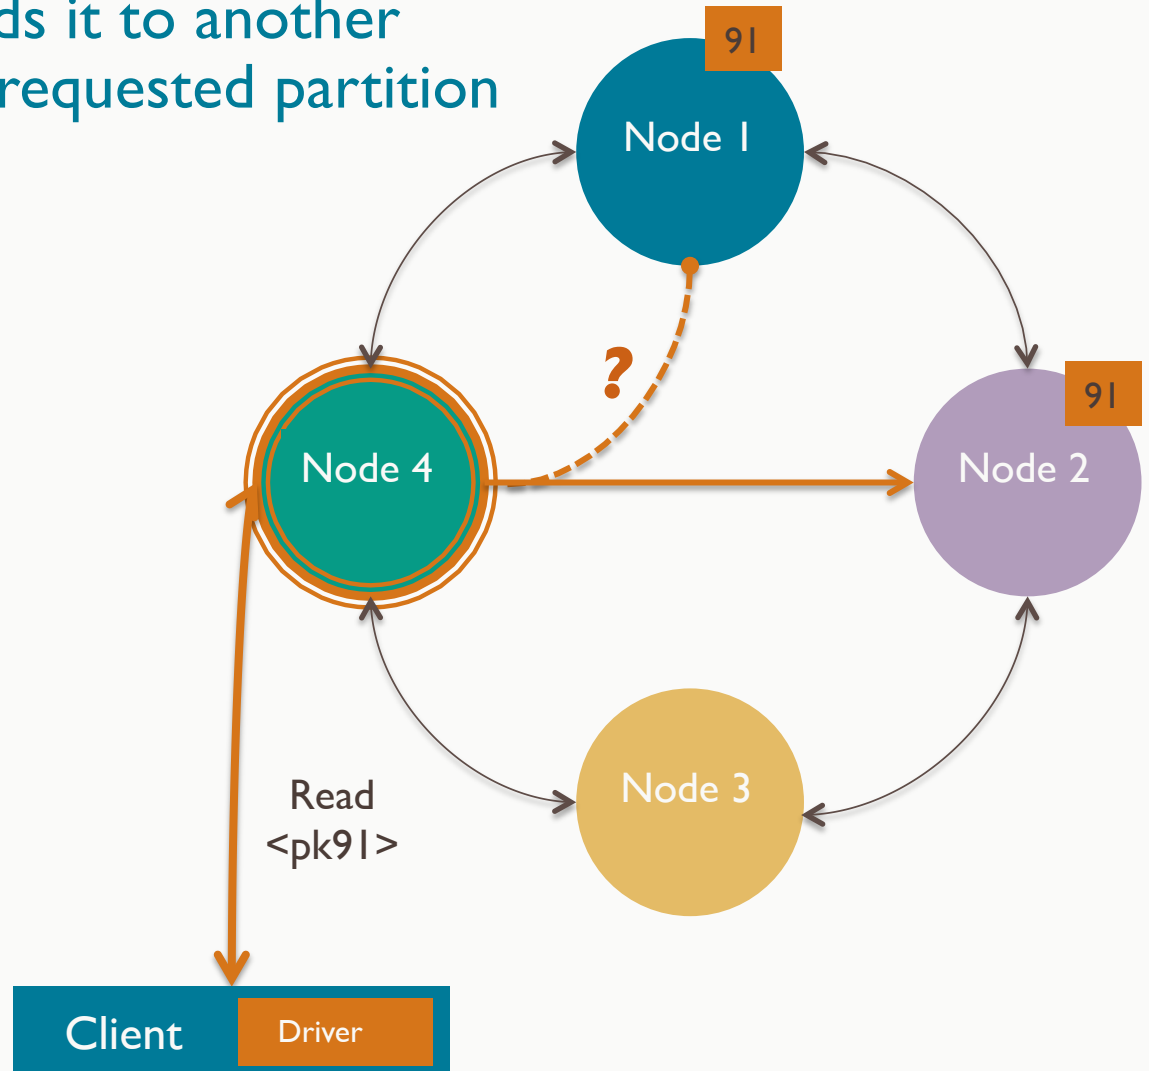
How does the read path flow among nodes?

- Cassandra returns the most recent record among the nodes read for a given request
 - Consistency Level
 - sets how many nodes will be read for a given request
 - may vary by request
- 



What are eager retries?

- If a node is slow responding to a request, the coordinator forwards it to another holding a replica of the requested partition
 - New feature in 2.0+
 - Only relevant if $RF > 1$



What are the key components of the read path?

- Each node implements in-memory structures for each CQL table
 - **MemTable** – in-memory table serves data as part of the *merge* process
 - **Row Cache** – in-memory cache stores recently read rows (optional)
 - **Bloom Filters** – reports if a *partition key* may be in its corresponding *SSTable*
 - **Key Caches** – maps recently read *partition keys* to specific *SSTable* offsets
 - **Partition Summaries** – Sampling from *partition index*
- Each node implements these on disk for each CQL table
 - **Partition Indexes** – Sorted *partition keys* mapped to their *SSTable* offsets
 - **SSTables** – static files periodically flushed from a *MemTable*
- **Merge** – unless served from the *row cache*, a read uses a *partition key* to locate, merge, and return values from a *MemTable* and any related *SSTable* storing values for that key

How does the read path flow on *each* node?

Row cache hit

Off Heap On Heap

Coordinator

Read
<pk7>

Hit

Row Cache (optional)
pk1, pk2, **pk7**

pk7	first: Elizabeth	last: Blue	level: 42
------------	-------------------------	-------------------	------------------

MemTable (e.g., player)

...
pk7	level: 42 timestamp 1114

Node memory
Node file system

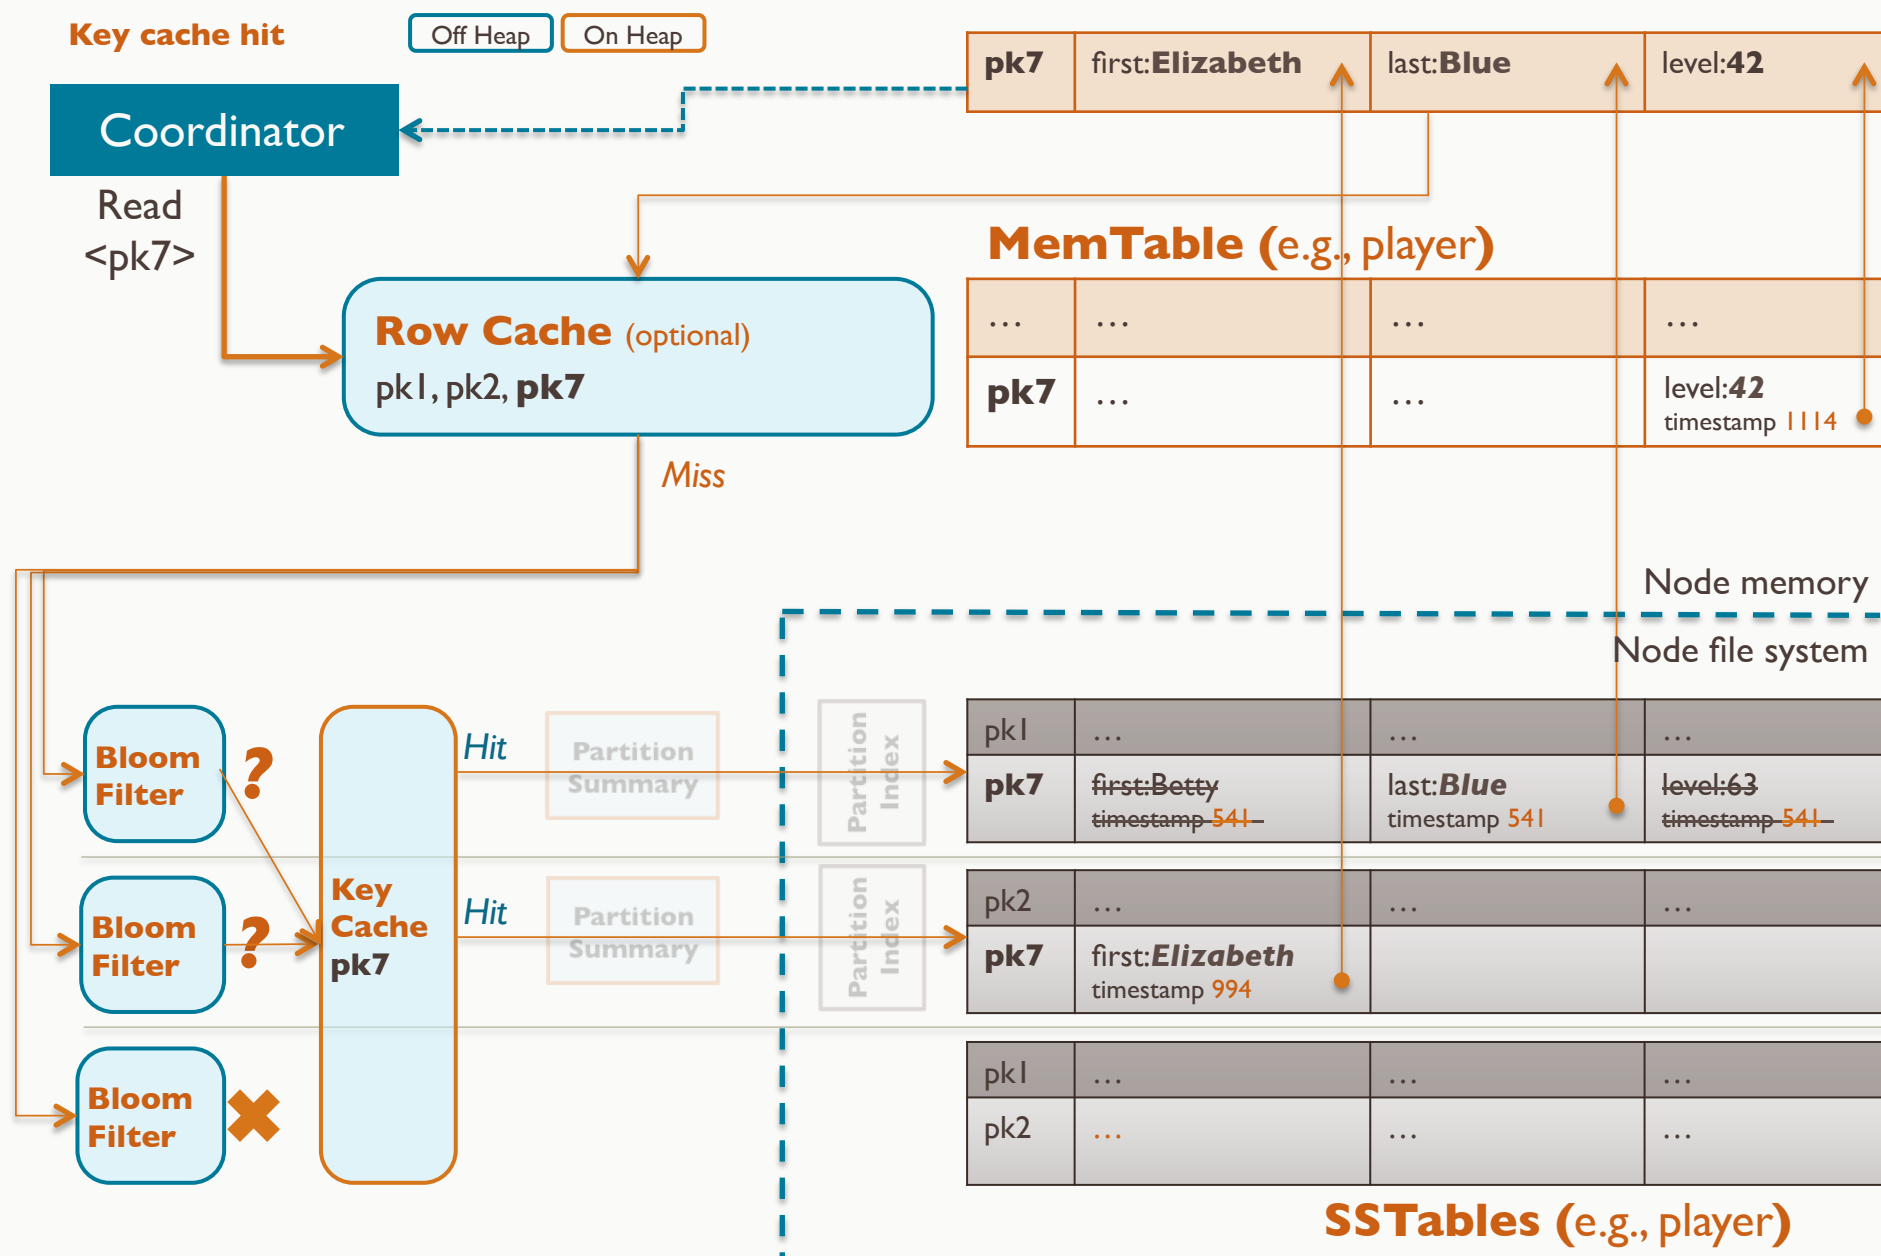
pk1
pk7	first: Betty timestamp 541	last: Blue timestamp 541	level: 63 timestamp 541

pk2
pk7	first: Elizabeth timestamp 994		

pk1
pk2

SSTables (e.g., player)

How does the read path flow on *each* node?

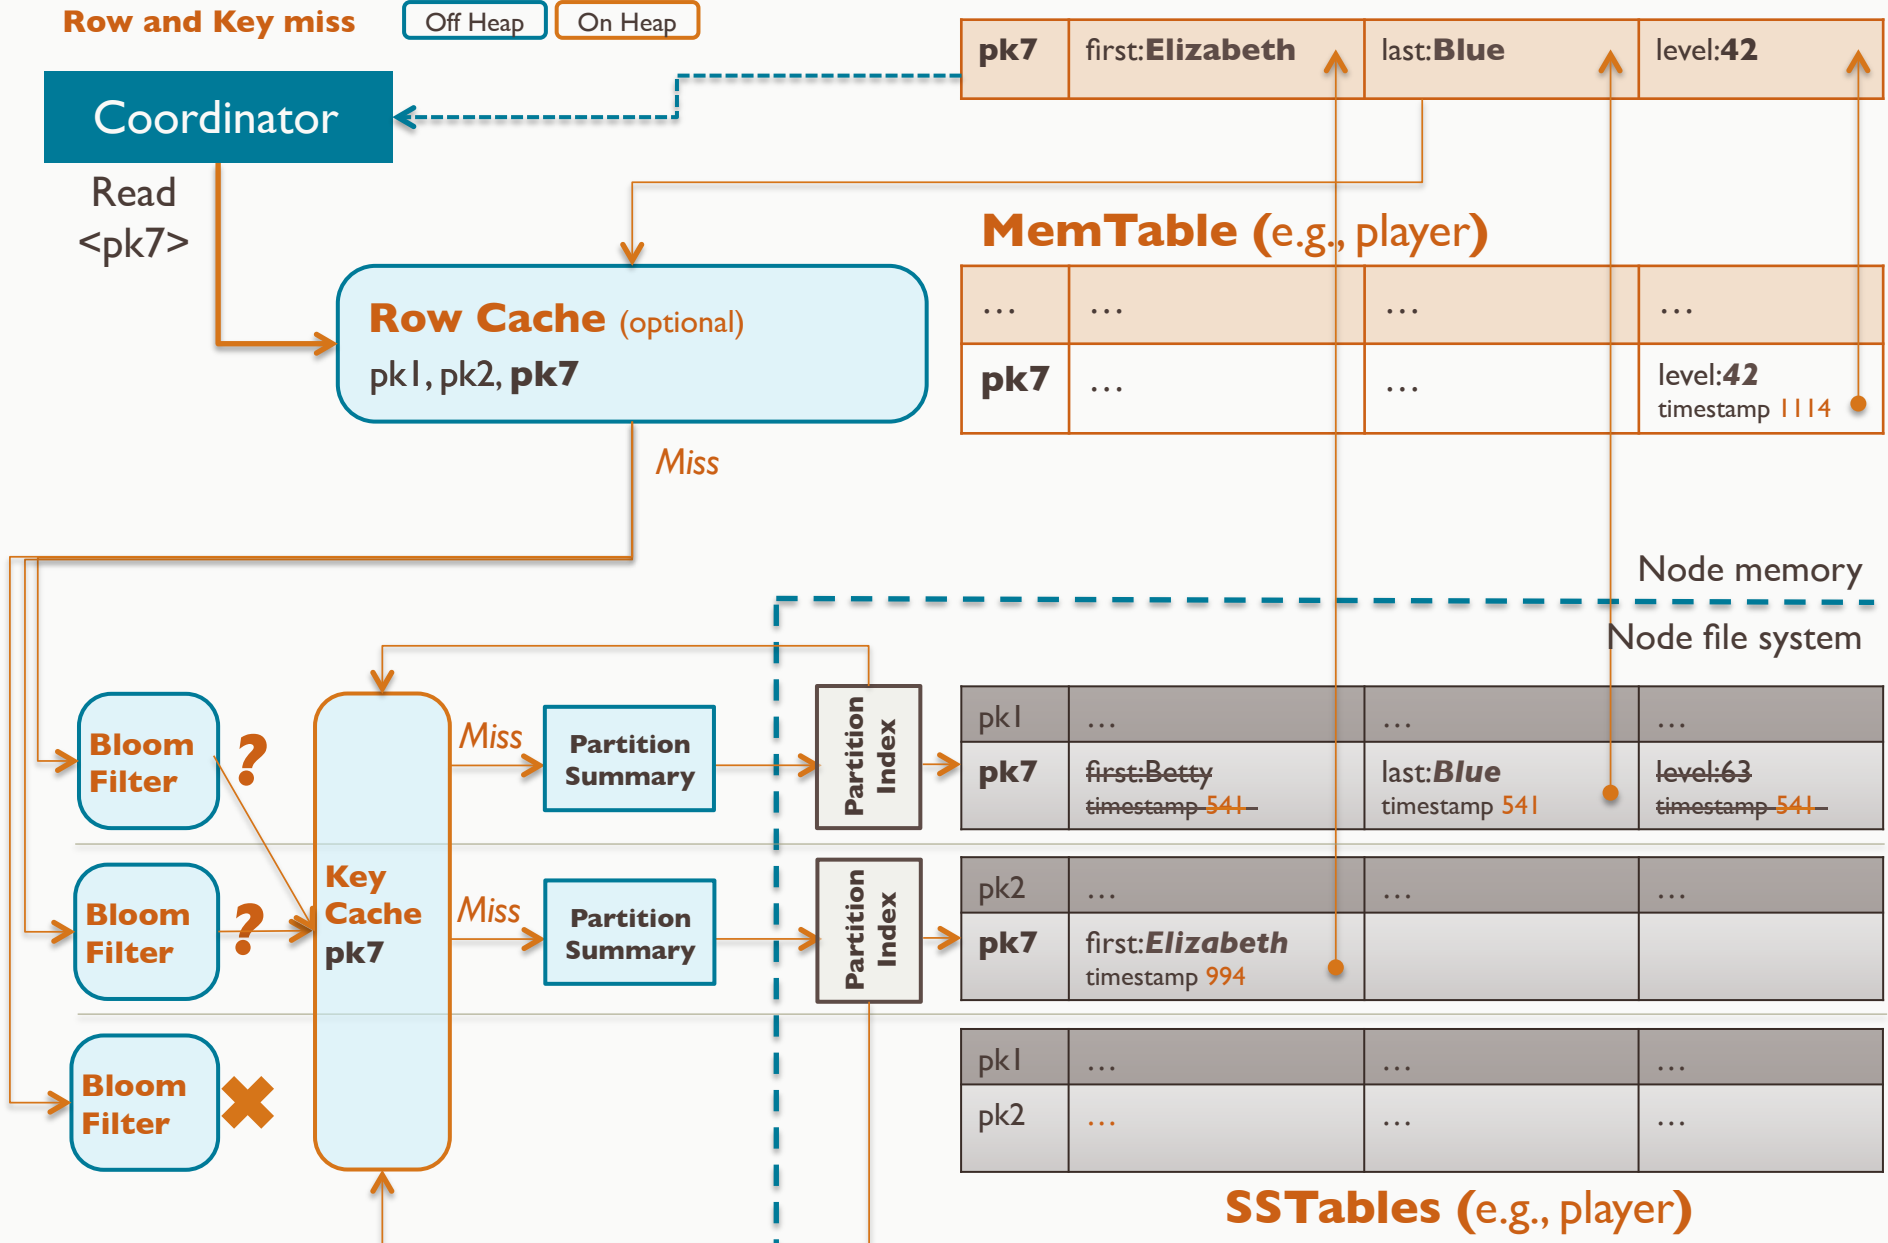


How does the read path flow on *each* node?

Row and Key miss

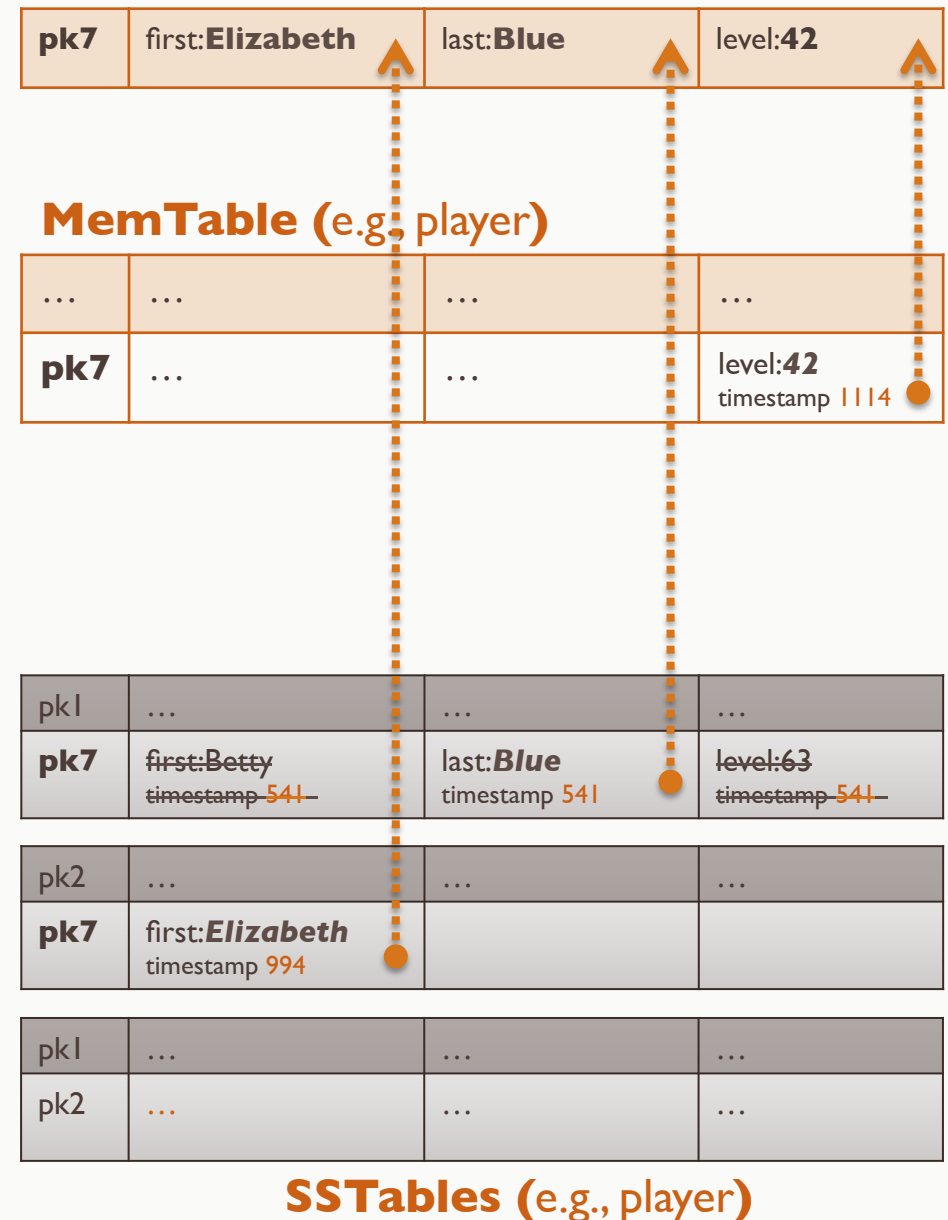
Off Heap

On Heap



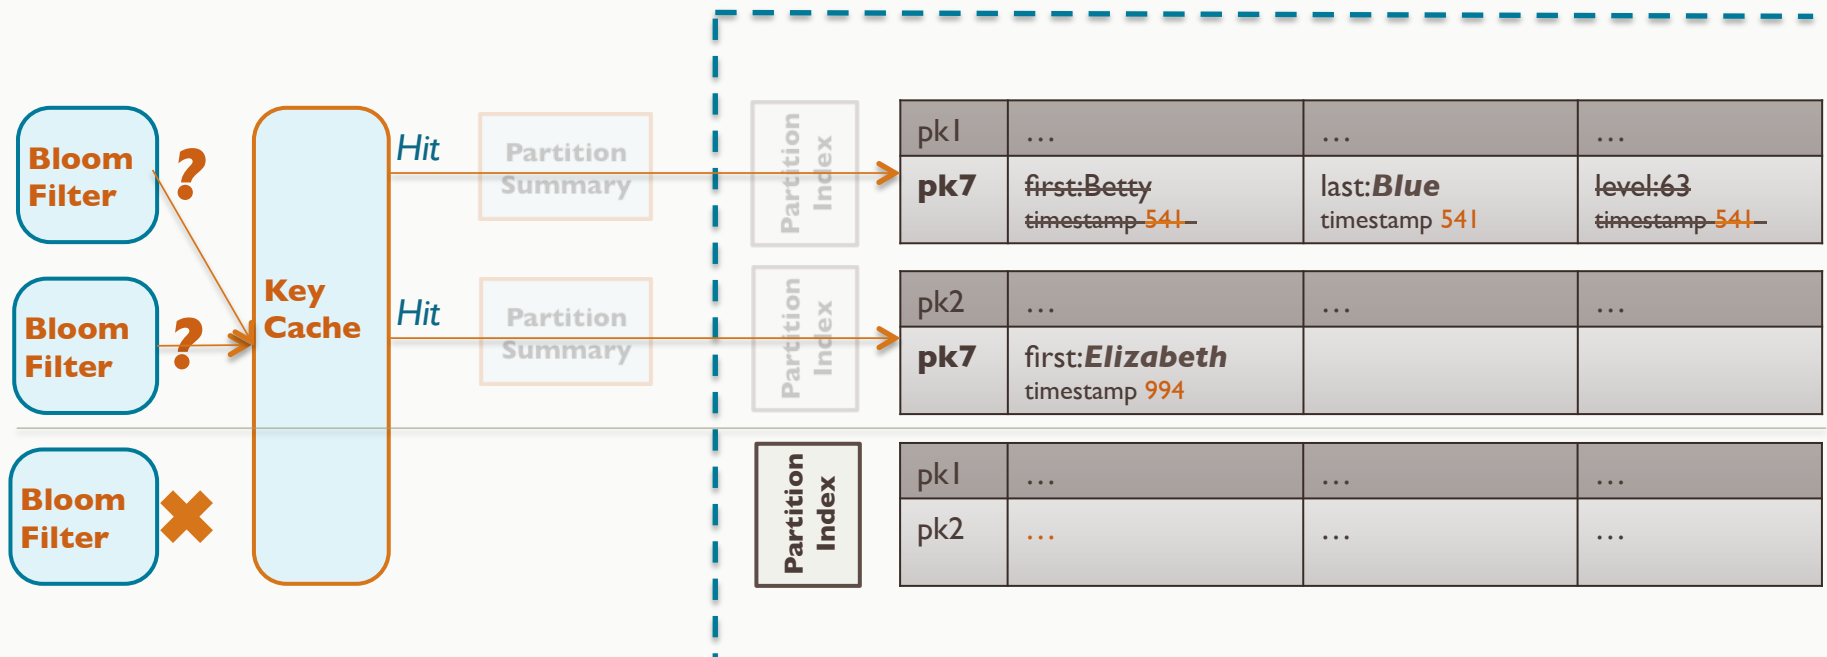
How is a MemTable and its SSTables used during a read?

- Both a *MemTable* and its recent *SSTables* are checked when reading for a partition key
 - the most current column values are combined to form the result



What is a Bloom filter and how does it optimize a read?

- A probabilistic data structure testing if a key may be in a SSTable
 - each SSTable has a Bloom filter on disk, used from off-heap memory
 - false positives are possible, false negatives are not
 - larger tables have a higher possibility of false positives
 - 1gb to 2gb per billion partitions in a SSTable
- Eliminates seeking a partition key in any SSTable without it



What is the bloom_filter_fp_chance table setting?

- Controls the percentage chance of false positive results from the Bloom filters for SSTables flushed for a specified table
- Values range from 0.0 to 1.0
 - 0.0 no false positives, greatest memory use
 - 0.1 maximum recommended setting, diminishing returns if higher
 - 1.0 Bloom filtering disabled for this table
- Default setting depends on compaction strategy
 - 0.01 Size-tiered compaction (STC)
 - 0.1 Leveled compaction (LCS)

```
ALTER TABLE player  
WITH bloom_filter_fp_chance = 0.1;
```

Exercise I: Working with Bloom filters



What is the row cache and how is it configured?



- The *merged* row(s) for a *partition* key is saved in off-heap memory
- Row caching is enabled in CQL with the *caching* and *rows_per_partition* properties
 - **ALL** – cache all rows for a partition key
 - **n** – cache the first *n* rows for a partition key
 - **NONE** – (default) disable row caching for this table

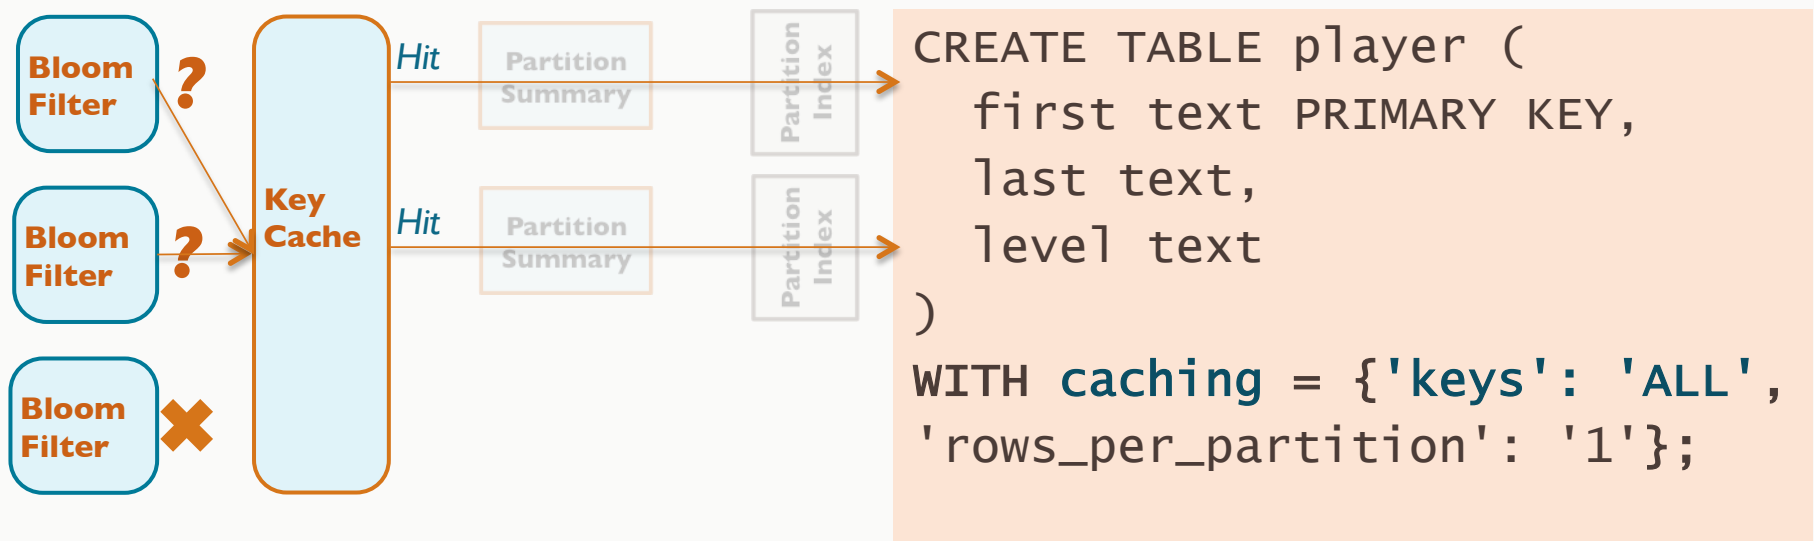
```
CREATE TABLE player (
    first text PRIMARY KEY,
    last text,
    level text
)
WITH caching = {'keys': 'ALL',
'rows_per_partition': '1'};
```

What is the row cache and how is it configured?

- Caches can periodically save to disk improving node restart speed
- Row cache size and save period are set globally for all tables on a node in *cassandra.yaml*
 - **row_cache_size_in_mb** – maximum row cache size, set to 0 to disable
 - **row_cache_save_period** – periodicity in seconds at which row cache should be saved to the *saved_caches_directory*, improves cache usage following a node restart, set to 0 to disable row cache saving
 - **row_cache_keys_to_save** – max number of cached rows to save each period, if disabled all cached rows are saved
 - **saved_caches_directory** – location to save row, key, and counter caches
 - default: `/var/lib/cassandra/saved_caches`

What is the key cache and how is it configured?

- Key caching saves a *partition key* and its *offset position(s)* in the *SSTable(s)* for a *MemTable*
 - one key cache entry for each *SSTable* holding a *replica* of this partition
 - reduces a read to a single seek per recent replica
- Key caching is enabled in CQL with the *caching* and *keys* properties
 - **ALL** – (default) enable key caching for this table
 - **NONE** – disable key caching for this table



What is the key cache and how is it configured?

- Caches can periodically save to disk, to improve node restart speed
- Key cache size and save period are set globally for all tables on a node in *cassandra.yaml*
 - **key_cache_size_in_mb** – maximum key cache size, set to 0 to disable
 - default: 5% of available heap or 100mb, whichever is smaller
 - **key_cache_save_period** – periodicity in seconds at which key cache should be saved to the *saved_caches_directory*, improves cache usage following a node restart, set to 0 to disable key cache saving
 - **key_cache_keys_to_save** – max number of cached keys to save each period, if disabled all cached keys are saved
 - **saved_caches_directory** – location to save row, key, and counter caches
 - default: /var/lib/cassandra/saved_caches
- Enabling key caching is commonly termed "pre-heating"

What is the counter cache and how is it configured?

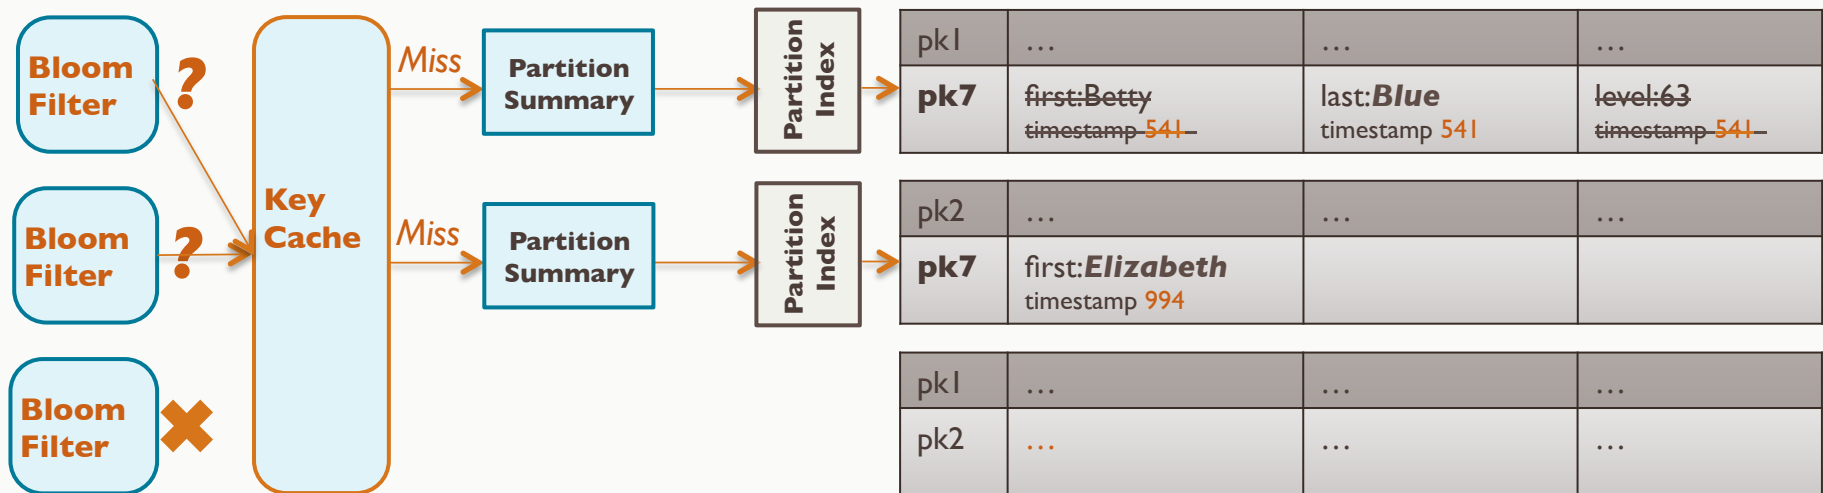
- Counter caching saves the clock and count of a counter in memory
 - helps reduce lock contention for the read-before-write counter updates
 - only the local tuple for node is saved in the counter cache
- Counter cache size and save period are set globally for all counter tables on a node in *cassandra.yaml*
 - **counter_cache_size_in_mb** – maximum counter cache size, set to 0 to disable
 - default: 2.5% of available heap or 50mb, whichever is smaller
 - **counter_cache_save_period** – periodicity in seconds at which the counter cache should be saved to the *saved_caches_directory*, improves cache usage following a node restart, set to 0 to disable counter cache saving
 - **counter_cache_keys_to_save** – max number of cached keys to save each period, if disabled all cached keys are saved
 - **saved_caches_directory** – location to save row, key, and counter caches
 - default: /var/lib/cassandra/saved_caches

Exercise 2: Enable the key cache



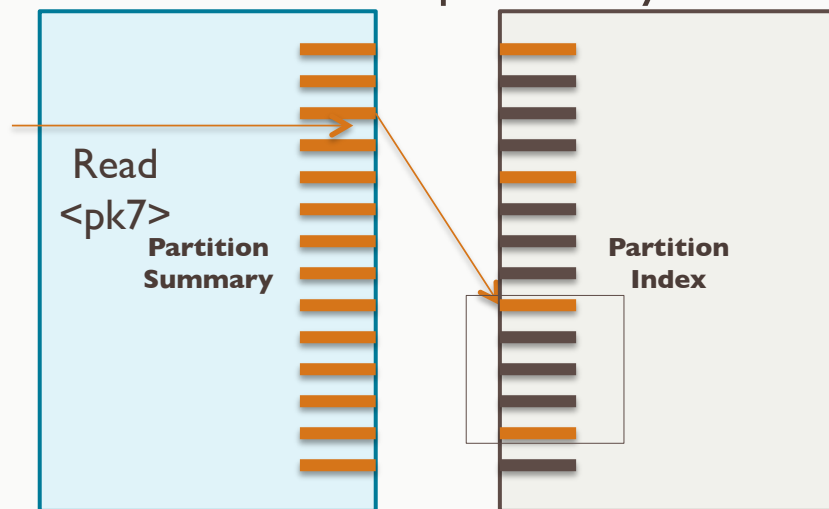
What are partition summaries and how are they used?

- If a partition key's location is not in the *key cache*, the read must seek the requested partition on disk
- The *partition summary* in an in-memory sampling from a *partition index*, used to locate a key's approximate location in the full index
 - default sample ratio is 1 per 128 partition keys in the index
 - configured with the table property `min_index_interval` (default: 128) and `max_index_interval` (default: 2048)
 - held in off-heap memory



What are partition summaries and how are they used?

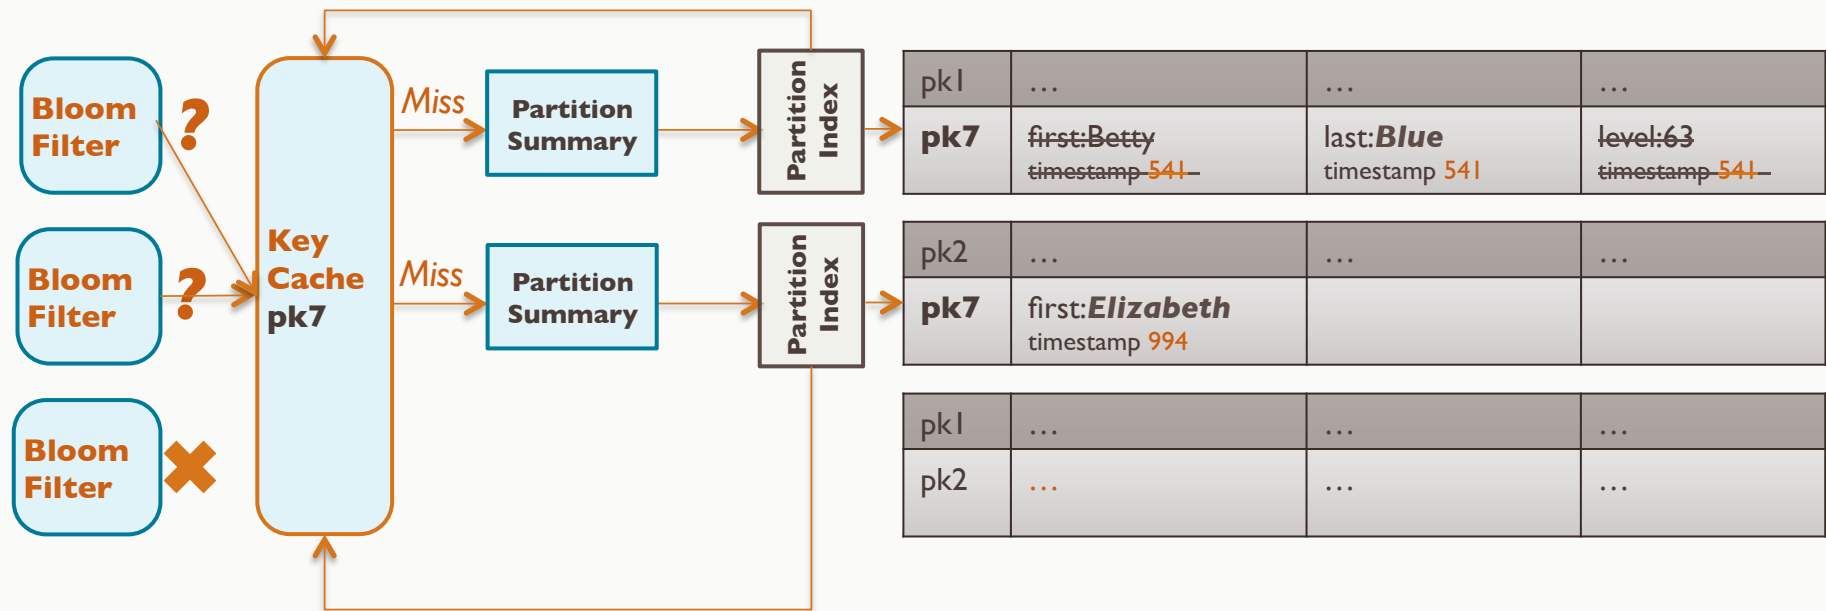
- If a partition key's location is not in the *key cache*, the read must seek the requested partition on disk
- The *partition summary* is an in-memory sampling from a *partition index*, used to locate a key's approximate location in the full index
 - default sample ratio is 1 per 128 partition keys in the index
 - configured with the table property `min_index_interval` (default: 128) and `max_index_interval` (default: 2048)
 - held in off-heap memory



```
CREATE TABLE player (  
    first text PRIMARY KEY,  
    last text,  
    level text  
)  
  
WITH min_index_interval = 256  
AND max_index_interval = 2048;
```


What are partition indexes and how are they used?

- The *partition index* of each SSTable provides the physical offset locations for each of its partitions, sorted by *partition key*
- Starting with the approximate location from the *partition summary*, the *partition index* is read to find the physical position of a partition
 - Once found, the location of this partition key is added to the *key cache*



Learning Objectives

- Understand how data is read from the storage engine
- **Read data from Cassandra**

How do you execute CQL queries in cqlsh?

- As learned earlier, *cqlsh* enables command line CQL execution

```

dstraining@DST: /home/cassandra
dstraining@DST:/home/cassandra$ bin/cqlsh
Connected to Test Cluster at localhost:9160.
[cqlsh 4.1.1 | Cassandra 2.0.5 | CQL spec 3.1.1 | Thrift protocol 19.39.0]
Use HELP for help.
cqlsh> DESCRIBE KEYSPACES;
system music system_traces demo

cqlsh> USE music;
cqlsh:music> SELECT *
... FROM performer
... LIMIT 5;

name | born | country | died | founded | style | type
-----+-----+-----+-----+-----+-----+-----
Sheryl Crow | 1962 | United States | null | null | Rock | artist
Black Bottle Scotch Whisky Pipe Band | null | Scotland | null | 1989 | Pipe and Drum | band
Bellefire | null | Ireland | null | null | Unknown | band
Dia DiCristino | null | United States | null | null | Unknown | artist
Pat Green | null | United States | null | null | Unknown | artist

(5 rows)
cqlsh:music>
  
```

Note, cqlsh and CQL are taught in detail in Module 4 – Introducing the Cassandra Data Model and CQL

How do you examine data storage using CLI?

- The *cassandra-cli* utility is
 - useful for examining and learning the internal storage engine structure
 - deprecated and less functional than CQL and *cqlsh*, which are fully backwards-compatible with column families and data created using Thrift

use	create	set	get	list
limit	help	assume	quit	

```
dstraining@DST: /home/cassandra
dstraining@DST:/home/cassandra$ bin/cassandra-cli
Connected to: "Test Cluster" on 127.0.0.1/9160
Welcome to Cassandra CLI version 2.0.5

[default@unknown] USE music;
Authenticated to keyspace: music
[default@music] LIST performer LIMIT 1;
Using default cell limit of 100
-----
RowKey: Sheryl Crow
=> (name=, value=, timestamp=1394065306886002)
=> (name=born, value=000007aa, timestamp=1394065306886002)
=> (name=country, value=556e6974656420537461746573, timestamp=1394065306886002)
=> (name=style, value=526f636b, timestamp=1394065306886002)
=> (name=type, value=617274697374, timestamp=1394065306886002)
```

Demo 3: Use the CLI to examine data storage



How is CQL tracing enabled and used?

- The CQL TRACING command enables and disables request tracing
 - results displayed and saved to *sessions* and *events* in *system_traces* keyspace

```
cqlsh:musicdb> TRACING ON;
Now tracing requests.
cqlsh:musicdb> SELECT * FROM performer WHERE name = 'Sheryl Crow';
```

name	born	country	died	founded	style	type
Sheryl Crow	1962	United States	null	null	Rock	artist

```
(1 rows)
```

Tracing session: 291832f0-3e13-11e4-898b-17914c10dbe5

activity	timestamp	source	source_elapsed
Execute CQL3 query	2014-09-16 19:34:34.271000	127.0.0.1	0
Parsing SELECT * FROM performer WHERE name = 'Sheryl Crow' LIMIT 10000;	[SharedPool-Worker-1] 2014-09-16 19:34:34.271000	127.0.0.1	86
Preparing statement	[SharedPool-Worker-1] 2014-09-16 19:34:34.271000	127.0.0.1	180
Executing single-partition query on performer	[SharedPool-Worker-2] 2014-09-16 19:34:34.271000	127.0.0.1	595
Acquiring sstable references	[SharedPool-Worker-2] 2014-09-16 19:34:34.271000	127.0.0.1	614
Merging memtable tombstones	[SharedPool-Worker-2] 2014-09-16 19:34:34.271000	127.0.0.1	656
Key cache hit for sstable 1	[SharedPool-Worker-2] 2014-09-16 19:34:34.272000	127.0.0.1	817
Seeking to partition beginning in data file	[SharedPool-Worker-2] 2014-09-16 19:34:34.272000	127.0.0.1	831
Skipped 0/1 non-slice-intersecting sstables, included 0 due to tombstones	[SharedPool-Worker-2] 2014-09-16 19:34:34.274000	127.0.0.1	2705
Merging data from memtables and 1 sstables	[SharedPool-Worker-2] 2014-09-16 19:34:34.274000	127.0.0.1	2729
Read 1 live and 2 tombstoned cells	[SharedPool-Worker-2] 2014-09-16 19:34:34.274000	127.0.0.1	2806
Request complete	2014-09-16 19:34:34.274413	127.0.0.1	3413

```
cqlsh:musicdb> TRACING OFF;
Disabled tracing.
```


Exercise 4: Read data and examine its tracing output



How do you obtain performance data using cfstats?

- *nodetool cfstats* command provides statistics for a specified table ("column family"), including
 - read latency
 - write latency
 - SSTable count
 - space used

```
dstraining@DST:/home/cassandra$ bin/nodetool cfstats musicdb.performer
-----
Keyspace: musicdb
  Read Count: 4
  Read Latency: 0.20725 ms.
  Write Count: 5537
  Write Latency: 0.03769297453494672 ms.
  Pending Flushes: 0
    Table: performer
      SSTable count: 1
      Space used (live), bytes: 550467
      Space used (total), bytes: 550467
      Space used by snapshots (total), bytes: 0
      SSTable Compression Ratio: 0.3156983447202369
      Memtable cell count: 0
      Memtable data size, bytes: 0
      Memtable switch count: 1
      Local read count: 4
      Local read latency: 0.208 ms
      Local write count: 5537
      Local write latency: 0.038 ms
      Pending flushes: 0
      Bloom filter false positives: 0
      Bloom filter false ratio: 0.00000
      Bloom filter space used, bytes: 6936
      Compacted partition minimum bytes: 30
      Compacted partition maximum bytes: 310
      Compacted partition mean bytes: 243
      Average live cells per slice (last five minutes): 1.0
      Average tombstones per slice (last five minutes): 2.0
```

```
bin/nodetool -h [host] -p [port] cfstats <keyspace>.<table>
```


Exercise 5: Use *cfstats* to measure performance



Summary

- If a node responds slowly to a request, the request is forwarded to another replica node
- The *row cache* is an optional mechanism to cache recently requested partitions
- Each SSTable has a *Bloom filter*, *partition summary*, and *partition index*
- A *Bloom filter* reduces disk seeks by ruling out SSTables which do not contain a partition
- The *key cache*, shared by all SSTables for a MemTable, caches the location of recently requested partition keys
- A *partition summary* is an evenly distributed in-memory sampling from a partition index, used to reduce index seek time
- A *partition index* provides the specific data file offset location for each partition key in an SSTable

Summary

- A *partition* key found in the *partition index* is added to the *key cache*
- Cassandra *merges* the most recent columns of data from a *MemTable*, and its *SSTables*, for a given request
- If *row cache* is in use, the merged *CQL* row for a partition key is updated when the requested row is returned
- Row and key caching is controlled using the *caching* table property
- *cqlsh* enables command line *CQL* queries and shell commands
- *cassandra-cli* enables command line Thrift API commands
- The *nodetool cfstats* command provides statistical information about a keyspace and table

Review Questions

- What benefit do Bloom filters provide to the read process?
- Is the partition summary read for partition keys in the key cache?
- What is the relationship between the partition summary and index?
- How many key caches are maintained for a MemTable?

